

Homework 2

600.482/682 Deep Learning

Spring 2020

February 14, 2020

Due Sun Feb. 23 11:59pm.

**Please submit a report (LaTeX generated PDF) and
the notebook as python file (file → download .py)
to Gradescope with entry code 9G83Y7
(submit the code as programming assignment)**

1. The goal of this problem is to minimize a function given a certain input using gradient descent by breaking down the overall function into smaller components via a computation graph. The function is defined as:

$$f(x_1, x_2, w_1, w_2) = \frac{1}{1 + e^{-(w_1 x_1 + w_2 x_2)}} + 0.5(w_1^2 + w_2^2).$$

- (a) Please calculate $\frac{\partial f}{\partial w_1}, \frac{\partial f}{\partial w_2}, \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}$.
 - (b) Start with the following initialization: $w_1 = 0.3, w_2 = -0.5, x_1 = 0.2, x_2 = 0.4$, draw the computation graph. Please use backpropagation as we did in class.
You can draw the graph on paper and insert a photo into your report.
The goal is for you to practice working with computation graphs. As a consequence, you must include the intermediate values during the forward and backward pass.
 - (c) Implement the above computation graph in the complimentary Colab Notebook using numpy. Use the values of (b) to initialize the weights and fix the input. Use a constant step size of 0.01. Plot the weight value w_1 and w_2 for 30 iterations in a single figure in the report.
2. The goal of this problem is to understand the classification ability of a neural network. Specifically, we consider the XOR problem. Go to the link in footnote¹ and answer the following questions. *Hint: hit reset the network right next to the run button after you change the architecture.*
 - (a) Can a linear classifier, without any hidden layers, solve the XOR problem?
 - (b) With one hidden layer and $\text{ReLU}(x) = \max(0, x)$, how many neurons in the hidden layer do you need to solve the XOR problem? Describe the training loss and estimated prediction accuracy when using 2, 3 and 4 neurons. Discuss the intuition of why a certain number of neurons is necessary to solve XOR.
 3. In this problem, we want to build a neural network from scratch using Numpy for a real-world problem. We consider the MNIST dataset (<http://yann.lecun.com/exdb/mnist/>), a hand-written digit classification dataset. Please follow the formula in the complimentary Colab Notebook. *Hint: Make sure you pass the loss and gradient check in the notebook.*

¹<https://playground.tensorflow.org/#activation=relu&batchSize=10&dataset=xor®Dataset=reg-plane&learningRate=0.01®ularizationRate=0&noise=0&networkShape=&seed=0.10699&showTestData=false&discretize=true&percTrainData=80&x=true&y=true&xTimesY=false&xSquared=false&ySquared=false&cosX=false&sinX=false&cosY=false&sinY=false&collectStats=false&problem=classification&initZero=false&hideText=false>

- (a) Implement the loss and gradient of a linear classifier (python function `linear_classifier_forward_and_backward`).
- (b) Implement the loss and gradient of a multilayer perceptron with one hidden layer and $\text{ReLU}(x) = \max(0, x)$ (python function `mlp_single_hidden_forward_and_backward`).
- (c) Implement the loss and gradient of a multilayer perceptron with two hidden layer, skip connection and $\text{ReLU}(x) = \max(0, x)$ (python function `mlp_two_hidden_forward_and_backward`).
- (d) Plot the development accuracy of each epoch of three models in a single figure using the following hyperparameters: the batch size is 50, the learning rate is 0.005 and the number of epochs is 20.
- (e) Try using other hyperparameters and select a set of best hyperparameters using **development accuracy**. Once you pick the best model and hyperparameters, include the development accuracy of each epoch into the above figure (make a new figure) and report the **test accuracy** of the selected model and hyperparameters.